

DistilBERT_detector

January 14, 2022

0.1 DistilBERT Detector to detect DGA domains.

0.1.1 Author: Abdulkarim Abdulkadir, s4840933

0.1.2 Load the libraries

We will load the libraries, and check if we are in the Google Colab environment to pip install ktrain and import the drive mount library. This is to make sure that if the notebook is run locally, it will not execute Google Colab environment commands.

```
[18]: import pandas as pd
import numpy as np
import os
import sys
from sklearn.utils import shuffle
from sklearn.model_selection import train_test_split

ENV_COLAB = 'google.colab' in sys.modules

if ENV_COLAB:
    ## install modules
    !pip install -q ktrain
    from google.colab import drive
    drive.mount('/content/drive', force_remount=True)

    ## print
    print('Environment: Google Colaboratory Pro+.')
    # import ktrain
```

Again we check which environment we are to correctly find the location of the data of our domains

```
[19]: if ENV_COLAB:
    dgaLocation = '/content/drive/MyDrive/research/DGA_domains/'
    benignDomains = '/content/drive/MyDrive/research/benign_domains/top-1m.csv'
else:
    dgaLocation = 'data/DGA_domains/'
    benignDomains = 'data/benign_domains/top-1m.csv'
```

We have a total amount of 19 different DGA types. Including the benign domain data, this will total 20 different types.

```
[20]: dgaDomains = [dga for dga in os.listdir(dgaLocation) if dga.endswith(r".csv")]
print("Total amount of DGA types: ", len(dgaDomains))
```

Total amount of DGA types: 37

0.2 Load the data into arrays

We will randomly select 110000 samples from the benign domains and 90000 from the dga domains. The ratio between the benign and dga domains will be 55:45. Thus we trimmed our data to a total of 200000 domains to use for training our BERT classifier

```
[21]: dataset = pd.DataFrame()
for i, dga in enumerate(dgaDomains):
    dgaDataFrame = pd.read_csv(dgaLocation + dga)
    dgaDataFrame.insert(1, 'type', dga.split(".")[0])
    dgaDataFrame.insert(2, 'class', 1)
    dataset = dataset.append(dgaDataFrame, ignore_index=True)
benignDataFrame = pd.read_csv(benignDomains)
benignDataFrame.insert(1, 'type', 'benign')
benignDataFrame.insert(2, 'class', 0)
dataset = dataset.append(benignDataFrame[:200000], ignore_index=True)
dataset = dataset.reset_index(drop=True)

[22]: print("Total amount of DGA domains: ", dataset['class'].value_counts()[1])
print("Total amount of benign domains: ", dataset['class'].value_counts()[0])
print("Total amount of domains: ", len(dataset))
if ENV_COLAB:
    dataset.to_csv('/content/drive/MyDrive/research/dataset', index=False)
else:
    dataset.to_csv('data/dataset', index=False)
```

Total amount of DGA domains: 184765
 Total amount of benign domains: 200000
 Total amount of domains: 384765

We will split our data into random train and test subsets. Our test size will be 25%. Our random_state that control the random number generated has to be given. Popular seeds are 42 or 0. We chose 42 for obvious reasons.

```
[32]: if ENV_COLAB:
    dataset = pd.read_csv('/content/drive/MyDrive/research/dataset')
else:
    dataset = pd.read_csv('data/dataset')
X = dataset['domain']

labels = dataset['class']
class_names = labels.unique()
```

```
x_train, x_test, y_train, y_test = train_test_split(dataset, labels, test_size=0.
→25, random_state=42)
```

Display the first 10 and last 10 data in our dataset.

```
[33]: display(dataset.head(10).append(dataset.tail(10)))
```

	domain	type	class
0	fliffmdhwrdb.org	locky	1
1	oqfyajoxgqnf.pw	locky	1
2	pnycpjwcw.pl	locky	1
3	bxicshg.info	locky	1
4	bamiitcnugeuexgt.work	locky	1
5	yuldhhytaiqgjrpfk.work	locky	1
6	irjhglaihb.xyz	locky	1
7	oxmblinqcstkj.xyz	locky	1
8	phudexoqds.xyz	locky	1
9	fqsfevovqshfcc.su	locky	1
384755	asumag.com	benign	0
384756	montessorinb.com	benign	0
384757	iremember.ru	benign	0
384758	hethongxulynuochai.com.vn	benign	0
384759	stevewonder.net	benign	0
384760	kenyan-post.com	benign	0
384761	adepem.com	benign	0
384762	gln.com	benign	0
384763	omahaoutdoors.com	benign	0
384764	talkhouse.com	benign	0

```
[36]: print("Size of training set: %s" % (len(x_train)))
print("Size of validation set: %s" % (len(x_test)))
```

Size of training set: 288573
Size of validation set: 96192

```
[37]: print(x_train.head(10))
```

	domain	type	class
284318	springeronline.com	benign	0
271248	indianss.org	benign	0
362013	pnsn.org	benign	0
274681	mtexpress.com	benign	0
106241	odonxllegulqj4t.com	shiotob	1
356934	chambre-agriculture.fr	benign	0
251804	s.coop	benign	0
207205	taishinbank.com.tw	benign	0
324992	neu.edu.vn	benign	0
46367	lyiemychun.com	fobber	1

We list all the text models that ktrain offers. For our research we will use the distilbert model. Which is a faster, smaller and distilled version of BERT.

```
[ ]: ktrain.text.print_text_classifiers()
```

```
fasttext: a fastText-like model [http://arxiv.org/pdf/1607.01759.pdf]
logreg: logistic regression using a trainable Embedding layer
nbsvm: NBSVM model [http://www.aclweb.org/anthology/P12-2018]
bigru: Bidirectional GRU with pretrained fasttext word vectors
[https://fasttext.cc/docs/en/crawl-vectors.html]
standard_gru: simple 2-layer GRU with randomly initialized embeddings
bert: Bidirectional Encoder Representations from Transformers (BERT) from
keras_bert [https://arxiv.org/abs/1810.04805]
distilbert: distilled, smaller, and faster BERT from Hugging Face transformers
[https://arxiv.org/abs/1910.01108]
```

```
[ ]: model_name = 'distilbert-base-uncased'
t = ktrain.text.Transformer(model_name, class_names=labels.unique(),
                             maxlen=350)
```

```
Downloading: 0%|          | 0.00/483 [00:00<?, ?B/s]
```

Naming our pre-process train and validation dataset respectively.

```
[ ]: train = t.preprocess_train(x_train.tolist(), y_train.to_list())

val = t.preprocess_test(x_test.tolist(), y_test.to_list())
model = t.get_classifier()
```

```
preprocessing train...
```

```
language: en
```

```
train sequence lengths:
```

```
mean : 1
```

```
95percentile : 1
```

```
99percentile : 1
```

```
Downloading: 0%|          | 0.00/232k [00:00<?, ?B/s]
```

```
Downloading: 0%|          | 0.00/466k [00:00<?, ?B/s]
```

```
Downloading: 0%|          | 0.00/28.0 [00:00<?, ?B/s]
```

```
<IPython.core.display.HTML object>
```

```
Is Multi-Label? False
```

```
preprocessing test...
```

```
language: en
```

```
test sequence lengths:
```

```
mean : 1
```

```
95percentile : 1
```

```
99percentile : 1
```

<IPython.core.display.HTML object>

Downloading: 0%| | 0.00/363M [00:00<?, ?B/s]

We will find a good learning rate using the learning rate range test to provide valuable information about an optimal learning rate. To point has to be chosen at which the loss starts descending and the point at which the loss stops descending or becomes ragged. For BERT and DistilBERT models the typical learning rate is between $5e-5$ and $2e-5$.

```
[ ]: learner = ktrain.get_learner(model,
                                train_data=train,
                                val_data=val,
                                batch_size=6)
```

```
[ ]: learner.lr_find(max_epochs=4)
      learner.lr_plot()
```

simulating training for different learning rates... this may take a few moments...

Epoch 1/4


25000/25000 [=====] - 3781s 150ms/step - loss: 0.1452 - accuracy: 0.9468

Epoch 2/4

25000/25000 [=====] - 2056s 81ms/step - loss: 0.0797 - accuracy: 0.9664

done.

Please invoke the `Learner.lr_plot()` method to visually inspect the loss plot to help identify the maximal learning rate associated with falling loss.



Based on the plot above we choose $3e-5$ as our learning rate. We will fit a model following the 1cycle policy.

```
[ ]: learner.fit_onecycle(3e-5, 4)
```

begin training using one cycle policy with max lr of 3e-05...

Epoch 1/4

25000/25000 [=====] - 3861s 153ms/step - loss: 0.0389 - accuracy: 0.9874 - val_loss: 0.0181 - val_accuracy: 0.9944

Epoch 2/4

25000/25000 [=====] - 3861s 154ms/step - loss: 0.0163 - accuracy: 0.9956 - val_loss: 0.0292 - val_accuracy: 0.9947

Epoch 3/4

25000/25000 [=====] - 3869s 154ms/step - loss: 0.0091 - accuracy: 0.9976 - val_loss: 0.0166 - val_accuracy: 0.9965

Epoch 4/4

25000/25000 [=====] - 3888s 155ms/step - loss: 0.0027 - accuracy: 0.9994 - val_loss: 0.0147 - val_accuracy: 0.9972

```
[ ]: <keras.callbacks.History at 0x7f3d7d5b2110>
```

Save the learned model to location, to reuse the model without having to learn our dataset again.

```
[ ]: predictor = ktrain.get_predictor(learner.model, preproc=t)
      if ENV_COLAB:
          predictor.save('/content/drive/MyDrive/research/model/')
      else:
          predictor.save('model/')
```

View observations in learner with top losses in validation dataset.

```
[ ]: learner.view_top_losses(preproc=t, n=1, val_data=None)
```

```
-----
id:29028 | loss:12.09 | true:1 | pred:0)
```


We will cross-check the model with the validation(test) data, we can use the validate method assigned to the learner instance of the model. The results will be displayed in precision, recall and f1 score.

```
[ ]: learner.validate(val_data=val, class_names=t.get_classes())
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	27737
1	1.00	1.00	1.00	22263
accuracy			1.00	50000
macro avg	1.00	1.00	1.00	50000
weighted avg	1.00	1.00	1.00	50000

```
[ ]: array([[27674, 63],
           [ 77, 22186]])
```

```
[ ]: learner.plot()
```



```
[ ]: valid_preds = learner.predict()
len(valid_preds), dataset.shape, valid_preds[:5]
```

```
[ ]: (50000, (200000, 3), array([[6.0326261e-06, 9.9999392e-01],
[9.9999654e-01, 3.3981382e-06],
[9.9999535e-01, 4.6324981e-06],
[9.9999630e-01, 3.6541069e-06],
[1.6295456e-04, 9.9983704e-01]]), dtype=float32))
```

0.2.1 Model prediction on validation data

Load the saved predictor model to predict on the validation data.

```
[ ]: if ENV_COLAB:
    predictor = ktrain.load_predictor('/content/drive/MyDrive/research/model/')
else:
    predictor = ktrain.load_predictor('model/')
learner = ktrain.get_learner(predictor.model, train_data = train, val_data =
    ↪ val, batch_size = 12)
```

```
[ ]: predictor.get_classes()
```

```
[ ]: [0, 1]
```



```
[ ]: print(type(val))
# val = t.preprocess_test(x_test.tolist(), y_test.to_list())
x_fake = pd.Series(['hooghelandt.nl', 'radboud.nl', 'cryptojedi.org', 'https://
→ais.usvisa-info.com/', '000directory.com.ar', '01-telecharger.com', '1001tur.
→ru', '02022222222.com', 'ovenrenthighlightstablerefuse.com',
→'rowrepeatwakeassociatebox.com', 'iucyyekyksiewqo.org', 'owwxonkponu.co',
→'myeqiiookymyokqs.org'])
y_fake = pd.Series([1,1,1,1,1,1,1,1,0,0,0,0,0])
x_test = x_test.append(x_fake)
y_test = y_test.append(y_fake)

print(x_test.tail(13))
print(y_test.tail(13))
```

```
<class 'ktrain.text.preprocessor.TransformerDataset'>
0          hooghelandt.nl
1          radboud.nl
2          cryptojedi.org
3      https://ais.usvisa-info.com/
4          000directory.com.ar
5          01-telecharger.com
6          1001tur.ru
7          02022222222.com
8      ovenrenthighlightstablerefuse.com
9      rowrepeatwakeassociatebox.com
10         iucyyekyksiewqo.org
11         owwxonkponu.co
12         myeqiiookymyokqs.org
dtype: object
0      1
1      1
2      1
3      1
4      1
5      1
6      1
7      1
8      0
9      0
10     0
11     0
12     0
dtype: int64
```

```
[ ]: val = t.preprocess_test(x_fake.tolist(), y_fake.to_list())
learner.validate(val_data=val, class_names=t.get_classes())
```

preprocessing test...

language: en

test sequence lengths:

mean : 1

95percentile : 1

99percentile : 1

<IPython.core.display.HTML object>

	precision	recall	f1-score	support
0	0.00	0.00	0.00	5.0
1	0.00	0.00	0.00	8.0
accuracy			0.00	13.0
macro avg	0.00	0.00	0.00	13.0
weighted avg	0.00	0.00	0.00	13.0

```
[ ]: array([[0, 5],  
          [8, 0]])
```

Highlight the text of the validation sample to explain the prediction.

```
[ ]: from sklearn.metrics import accuracy_score  
pred=predictor.predict(x_test.to_list())  
acc=accuracy_score(y_test.to_list(),pred)
```

```
[ ]: print(acc)
```

0.9964227902236256

Predict on benign and dga domains that are not in our validation data or training data to check our classifier resiliance to new data.

```
[ ]: print("Benign domains: ")  
print(predictor.predict('hooghelandt.nl'))  
print(predictor.predict('radboud.nl'))  
print(predictor.predict('cryptojedi.org'))  
print(predictor.predict('https://ais.usvisa-info.com/'))  
print(predictor.predict('000directory.com.ar'))  
print(predictor.predict('01-telecharger.com'))  
print(predictor.predict('1001tur.ru'))  
print(predictor.predict('02022222222.com'))  
print("DGA domains: ")  
print(predictor.predict('ovenrenthighlightstablerefuse.com'))  
print(predictor.predict('rowrepeatwakeassociatebox.com'))  
print(predictor.predict('rowrepeatwakeassociatebox.com'))  
print(predictor.predict('iucyyekyuksiewqo.org'))  
print(predictor.predict('myeqiiookymyokqs.org'))
```

```
print(predictor.predict('owwxonkponu.co'))
```

Benign domains:

0
0
0
0
0
0
0
0
0

DGA domains:

1
1
1
1
1
1
1

```
[ ]: if ENV_COLAB:
    validation_data_location = '/content/drive/MyDrive/research/validation_data.
    ↳CSV'
else:
    validation_data_location = 'validation_data.csv'
validation_dataset = pd.read_csv(validation_data_location)
validation_dataset = validation_dataset.drop(labels='number',axis=1)
validation_dataset = validation_dataset.drop(labels=range(500000,1000000),axis=0)
validation_dataset = validation_dataset.
    ↳drop(labels=range(1500000,1800000),axis=0)
validation_dataset.reset_index()
validation_dataset = validation_dataset.replace(to_replace='legit', value=0)
validation_dataset = validation_dataset.replace(to_replace='conficker', value=1)
validation_dataset = validation_dataset.replace(to_replace='cryptolocker',
    ↳value=1)
validation_dataset = validation_dataset.replace(to_replace='zeus', value=1)
validation_dataset = validation_dataset.replace(to_replace='pushdo', value=1)
validation_dataset = validation_dataset.replace(to_replace='rovnix', value=1)
validation_dataset = validation_dataset.replace(to_replace='tinba', value=1)
validation_dataset = validation_dataset.replace(to_replace='matsnu', value=1)
validation_dataset = validation_dataset.replace(to_replace='ramdo', value=1)

print(validation_dataset)
for data in validation_dataset:
    print(data)
```

domain class

0	google.com	0
1	facebook.com	0
2	youtube.com	0
3	baidu.com	0
4	yahoo.com	0
...
1499995	byaaemigrationforthese.com	1
1499996	wethesenecessarypursuing.com	1
1499997	tyrantofonoverarmstrial.com	1
1499998	absolutestagepartsthey.com	1
1499999	themrepeatedsuchatof.com	1

[1000000 rows x 2 columns]

domain

class

```
[ ]: false = 0
      true = 0
      for idx, row in validation_dataset.iterrows():
          if predictor.predict(row['domain']) == row['class']:
              true = true + 1
          else:
              false = false + 1
```